

Linux kernel Fuzzing

Andrey Konovalov <andreyknvl@gmail.com>

ZeroNights 2016, Hardware Village
November 17th 2016

Agenda

- Userspace Sanitizers (ASan, TSan, MSan, UBSan)
- Kernel Sanitizers (KASAN, KTSAN, KMSAN)
- Kernel syscall fuzzers (Trinity, syzkaller)
- USB fuzzing: FaceDancer21, vUSBf

Userspace Sanitizers

- AddressSanitizer (ASan)
 - detects use-after-free and out-of-bounds
- ThreadSanitizer (TSan)
 - detects data races and deadlocks
- MemorySanitizer (MSan)
 - detects uninitialized memory uses
- UndefinedBehaviorSanitizer (UBSan)
 - detects undefined behaviors in C/C++

Kernel Sanitizers

- KASAN (use-after-free and out-of-bounds)
 - CONFIG_KASAN available upstream since 4.0
- KTSAN (data-races and deadlocks)
 - prototype available at <https://github.com/google/ktsan>
- KMSAN (uninitialized-memory-use)
 - in early prototype stage


KernelAddressSanitizer (KASAN)

Two parts

- Compiler module
 - Instruments memory accesses
- Runtime part
 - Bug detection algorithm


Shadow byte

Any aligned 8 bytes may have 9 states:
N good bytes and 8 - N bad ($0 \leq N \leq 8$)



Memory mapping

$$\text{Shadow} = (\text{Addr} \gg 3) + \text{Offset}$$



x86-64 memory layout

0000000000000000 - 00007fffffffffff (=47 bits) user space, different per mm
hole caused by [48:63] sign extension

ffff800000000000 - ffff87fffffffffff (=43 bits) guard hole, reserved for hypervisor

ffff880000000000 - ffffc7fffffffffff (=64 TB) direct mapping of all phys. memory

ffffc80000000000 - ffffc8fffffffffff (=40 bits) hole

ffffc90000000000 - ffffe8fffffffffff (=45 bits) vmalloc/ioremap space

ffffe90000000000 - ffffe9fffffffffff (=40 bits) hole

fffffea0000000000 - ffffeafffffffffffff (=40 bits) virtual memory map (1TB)

... unused hole ...

fffffec0000000000 - fffffc0000000000 (=44 bits) kasan shadow memory (16TB)

... unused hole ...

ffffff0000000000 - ffffff7fffffffffff (=39 bits) %esp fixup stacks

... unused hole ...

ffffffff80000000 - fffffffffa0000000 (=512 MB) kernel text mapping, from phys 0


fffffffafa0000000 - ffffffff5fffff (=1525 MB) module mapping space

ffffffffffff600000 - ffffffffdffffff (=8 MB) vsyscalls

ffffffffffffe00000 - ffffffff0000000 (=2 MB) unused hole

Compiler instrumentation: 8 byte access


```
*a = ...
```



```
char *shadow = (a >> 3) + Offset;  
if (*shadow)  
    ReportError(a);  
*a = ...
```

Instrumentation: N byte access (N = 1, 2, 4)

```
*a = ...
```



```
char *shadow = (a >> 3) + Offset;  
if (*shadow && *shadow < (a & 7) + N)  
    ReportError(a);  
*a = ...
```

Run-time module

- Maps shadow memory
- Adds redzones around slab objects
- Poisons/unpoisons shadow on kfree/kmalloc
- Ensures delayed reuse of slab objects
- Poisons global redzones on startup
- Collects stack traces for kmalloc/kfree
- Prints error reports

Fuzzing the kernel

- Kernel inputs:
 - system calls
 - network
 - USB
 - bluetooth
 -

Kernel system call fuzzers

- Trinity (<https://github.com/kernelslack/trinity>)
- syzkaller (<https://github.com/google/syzkaller>)

syzkaller

Existing system call fuzzers

Trinity in essence:

```
syscall(rand(), rand(), rand());
```

Knows argument types, so more like:

```
syscall(rand(), rand_fd(), rand_addr());
```

- Tend to find shallow bugs
- Frequently no reproducers

Coverage-guided fuzzing

(Think AFL or libFuzzer)

```
void TestOneInput(const char *data, int size) {  
    /* do something with data */  
}
```

Fuzzer invokes the function with different inputs

Code coverage guiding:

- Corpus of “interesting” inputs
- Mutate and execute inputs from corpus
- If inputs gives new coverage, add it to corpus

Coverage for the Linux kernel

- Available upstream with CONFIG_KCOV
- GCC pass that inserts a function call into every basic block
- kernel debugfs extension that collects and exposes coverage per-thread

```
if (...) {  
    ...  
}
```



```
__fuzz_coverage();  
if (...) {  
    __fuzz_coverage();  
    ...  
}  
__fuzz_coverage();
```

Syscall description

Declarative description of all syscalls:

```
open(file filename, flags flags[open_flags],  
      mode flags[open_mode]) fd
```

```
read(fd fd, buf buffer[out], count len[buf])
```

```
close(fd fd)
```

```
open_flags = O_RDONLY, O_WRONLY, O_RDWR, O_APPEND ...
```

Rich syscall description

Knows discriminated syscalls:

fcntl\$dupfd(fd fd, cmd **const**[F_DUPFD], arg fd) fd

fcntl\$getownex(fd fd, cmd **const**[F_GETOWN_EX], arg ptr[out, **f_owner_ex**])

Knows layout of structs:

```
f_owner_ex {  
    type    flags[f_owner_type, int32]  
    pid     pid  
}
```

Has unions:

```
tun_buffer [  
    pi  tun_pi  
    hdr virtio_net_hdr  
] [varlen]
```

Resources

```
resource fd_bpf_map[fd]
```

```
resource fd_bpf_prog[fd]
```

```
bpf$MAP_CREATE(cmd const[BPF_MAP_CREATE], ...) fd_bpf_map
```

```
bpf_map_lookup_arg {
```

```
    map      fd_bpf_map
```

```
    key      buffer[in]
```

```
    val      buffer[out]
```

```
}
```

Programs

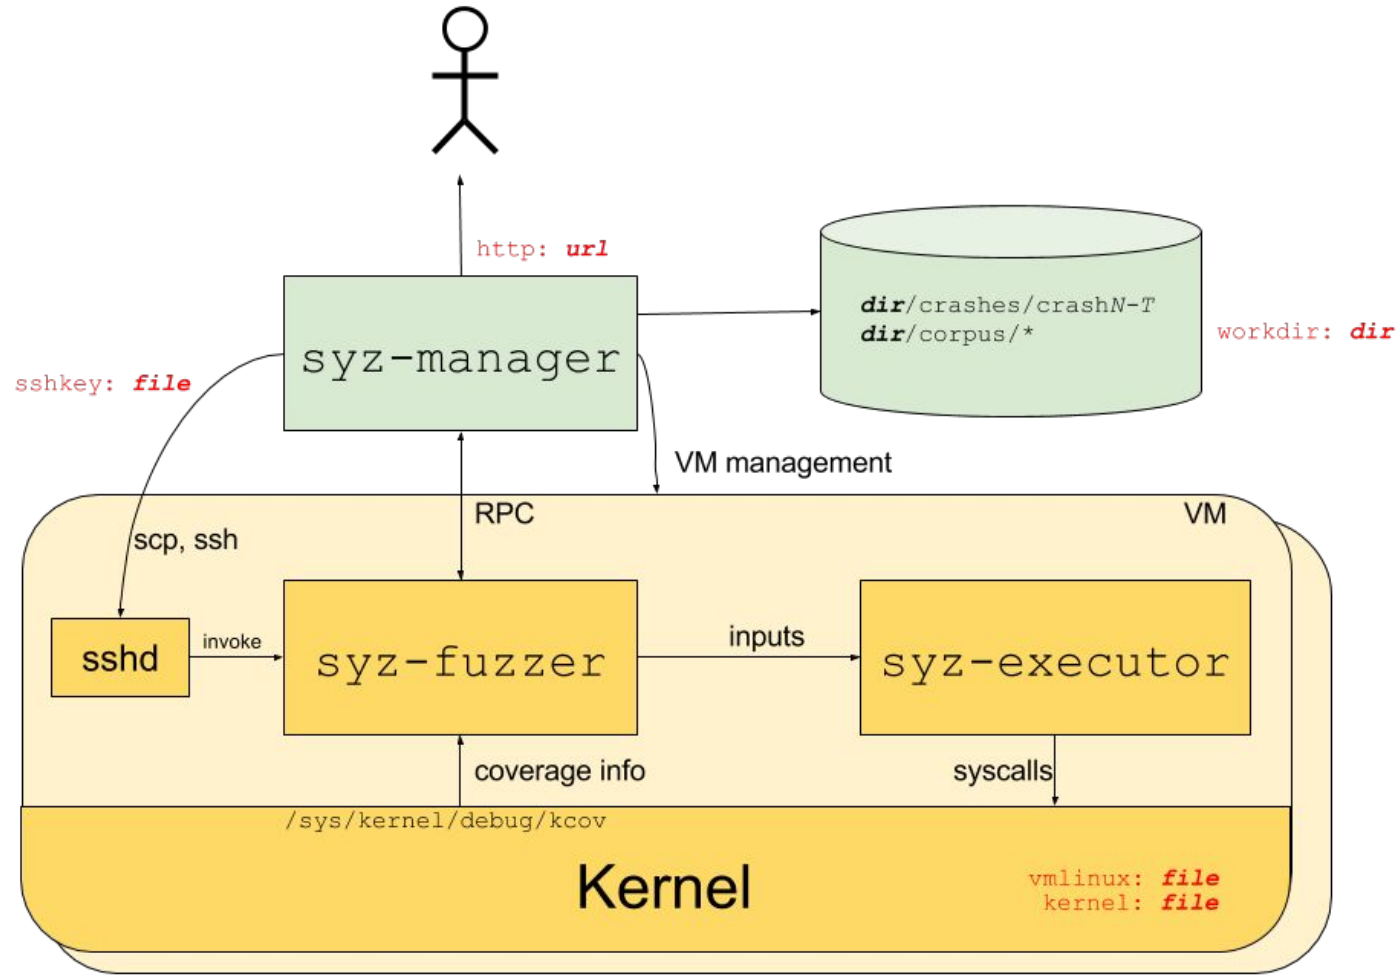
The description allows to generate and mutate "programs" in the following form:

```
mmap(&(0x7f0000000000), (0x1000), 0x3, 0x32, -1, 0)
r0 = open(&(0x7f0000000000)="./file0", 0x3, 0x9)
read(r0, &(0x7f0000000000), 42)
close(r0)
```

Algorithm

1. Start with empty corpus of programs
2. Generate a new program, or choose an existing program from corpus and mutate it (know argument types!)
3. Interpret the program, collect coverage from every syscall independently
4. If a syscall covers code that wasn't covered by this syscall previously, minimize program and add to corpus
5. Goto 1

Overview



External Stimulus

Systems calls and external stimulus in the same program:

```
listen(r0)
emit_ethernet(syn)
emit_ethernet(ack)
r1 = accept(r0)
emit_ethernet(data)
read(r1)
emit_ethernet(rst)
```

Work in progress; also applicable to USB, ...

Other Linux kernel fuzzers

- <https://github.com/oracle/kernel-fuzzing>
- <https://github.com/nccgroup/TriforceLinuxSyscallFuzzer>
- http://web.eece.maine.edu/~vweaver/projects/perf_events/fuzzer/
- <https://github.com/schumilo/vUSBf>


USB fuzzing

- Hardware: FaceDancer21
- In VM: vUSBf

(not BadUSB)

FaceDancer21

- “The purpose of this board is to allow USB devices to be written in host-side Python, so that one workstation can fuzz-test the USB device drivers of another host”
- <http://goodfet.sourceforge.net/hardware/facedancer21/>



FaceDancer21

- <https://github.com/travisgoodspeed/goodfet>
- <https://github.com/nccgroup/umap>
- <https://github.com/nccgroup/umap2>

vUSBf

- Virtual USB fuzzer
- QEMU + usbredir
- <https://github.com/schumilo/vUSBf>

CVE-2016-2384

- Double-free in usb-midi driver
- Found with vUSBf
- Confirmed and exploited with FaceDancer21
- <https://xairy.github.io/blog/2016/cve-2016-2384>

Questions?

<https://github.com/google/kasan/wiki>

<https://github.com/google/ktsan/wiki>

<https://github.com/google/syzkaller>

Andrey Konovalov, andreyknvl@gmail.com

kasan-dev@googlegroups.com

syzkaller@googlegroups.com